

Exercise 5 – AVL Tree

OOP 2014

May 7rd, 2014

1 Goals

1. Implementing a data structure learned in class (AVL Tree)
2. Experimenting with this data structure

2 Submission Details

- Submission Deadline: **Sunday, 18/5/2014, 23:55**
- This exercise will be done alone (**not in pairs**).
- You may use the *java.util.List* interface and two of its implementations – *java.util.LinkedList* and *java.util.ArrayList*. You may also use *java.util.Iterator*. You **may not** use any other class that you didn't write yourself.

3 Introduction

An AVL tree is a self-balancing binary search tree. AVL trees maintain the property that for each node, the difference between the heights of both its sub-trees is at most 1.¹ This is done by re-balancing the tree after each insertion and deletion operation. You recently learned about the theoretical background of AVL trees in the Data Structures (DaSt) course.²

In this exercise, you will implement an AVL tree of integers numbers based on the pseudo-code shown in the DaSt lecture.

4 API

You are required to submit a class *AvlTree* that implements the following API.

You are allowed (and encouraged!) to write and submit more class(es) as you see fit. You can also add methods of your own but make sure that you don't change the supplied API.

4.1 Package

Both *AvlTree* and any other class you implement in this exercise should be placed in the *oop.ex5.data_structures* package.

4.2 Methods

- ```
/**
 * Add a new node with key newValue into the tree.
 *
 * @param newValue new value to add to the tree.
 * @return false iff newValue already exist in the tree
 */
public boolean add(int newValue);
```

---

<sup>1</sup>The height of a tree is defined as the length of the longest downward path from the root to any of the leaves.

<sup>2</sup>See <http://moodle.cs.huji.ac.il/cs13/file.php/67109/tirgul8.pdf>.

- ```
/**
 * Does tree contain a given input value.
 *
 * @param val value to search for.
 * @return if val is found in the tree, return the depth of its node (where 0 is the root).
 * Otherwise – return -1.
 */
public int contains(int searchVal);
```
- ```
/**
 * Remove a node from the tree, if it exists.
 *
 * @param toDelete value to delete
 * @return true iff toDelete is found and deleted
 */
public boolean delete(int toDelete);
```
- ```
/**
 * @return number of nodes in the tree
 */
public int size();
```
- ```
/**
 * @return iterator to the Avl Tree. The returned iterator can pass over the tree nodes
 in ascending order.
 */
public Iterator<Integer> iterator();
```

### 4.3 Constructors

- ```
/**
 * A default constructor.
 */
public AvlTree();
```
- ```
/**
 * A data constructor –
 * a constructor that builds the tree by adding the elements in the input array one-by-one.
 * If the same value appears twice (or more) in the list, it is ignored.
 *
 * @param data values to add to tree
 */
public AvlTree(int[] data);
```
- ```
/**
 * A copy constructor –
 * a constructor that builds the tree a copy of an existing tree.
 *
 * @param tree an AvlTree
 */
public AvlTree(AvlTree tree);
```

4.4 Static Method

```
• /**
 * This method calculates the minimum number of nodes in an AVL tree of height h,
 *
 * @param h height of the tree (a non-negative number).
 * @return minimum number of nodes in the tree
 */
public static int findMinNodes(int h);
```

For example, calling `findMinNodes(3)`, as shown in Figure 1, should return 7

4.5 Comments

1. In the `add` method, in case the `newValue` already exist in the tree there is no need to add it twice.
2. The data constructor can get any array of integer number, that means you don't have any prior knowledge if the array is sorted or not.
3. Please **don't** paste code copied from the PDF file. (It may contain unwanted characters that will cause compilation problems).

5 Analyzing the AVL Tree

In Figure 1 you can see an AVL tree of height 3. This tree may seem unbalanced, however, this is in fact a valid AVL tree resulted from a specific insertion order. Notice that this example shows a tree of height 3 with minimal number nodes (if you delete any node, the height of the tree will be 2). Find a series of 12 numbers, such that when they are inserted into an empty AVL tree one-by-one, the result is a tree of height 4 (insertions only, no deletions). Write a full answer in the README file.

Hint: We suggest you start by figuring out in what order were the nodes inserted into the example tree, and then continue to create the larger one.

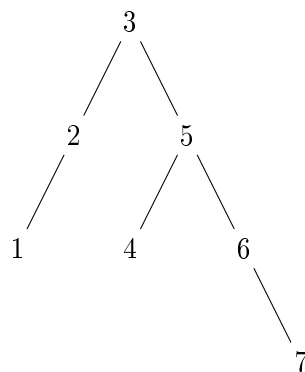


Figure 1: an AVL tree of height 3 with minimum number of nodes

6 Submission Requirements

6.1 README

Please address the following points in your README file:

1. Describe which class(es) (if any) you built in order to implement your AVL tree, other than *AvlTree*. The description should include the purpose of each class, its main methods and its interaction with the *AvlTree* class.
2. Describe your implementation of the methods *add()* and *delete()*. The description should include the general workflow in each of these methods. You should also indicate which helper functions you implemented for each of them, and which of these helper functions are shared by these two methods (if any).
3. Your README file should also answer the question from Section 5.

6.2 Automatic Testing

Our automatic tests will do the following:

1. Make sure your code compiles and contains a README file. This includes verifying that your *AvlTree* class follows the API defined in Section 4 (including the correct package names).
2. Test each of your methods on simple inputs, in order to test their basic functionality.
3. Test your tree on the four interesting AVL settings (i.e., the four rotations cases).
4. Test your code on complicated settings that include large inputs and multiple add/delete operations.
5. As always, we will run unseen tests which have 10% weight of the automatic testers.

In order to run our tests, we created a class called *TestEx5*, which uses an external *AvlTree* class. We will compile this class with each of the *AvlTree* classes submitted by the students. This class runs a set of tests on the *AvlTree* class it is compiled with. Each such test is defined by an input file, which creates an AVL tree and performs a set of operations on it. The format of the input files is described in Appendix A.

6.3 Submission Guidelines

You should submit a file named *ex5.jar* containing all the *.java* files of your program, as well as the README file. Please note the following:

- Files should be submitted along with their original packages.
- No *.class* files should be submitted.
- There is no need to submit any testers.
- Your program must compile without any errors or warnings.
- javadoc should accept your program and produce documentation.

You may use the following unix command to create the jar files:

```
jar -cvf ex5.jar README oop/ex5/data_structures/*.java
```

This command should be run from the main project directory (that is, where the *oop* directory is found).

7 Technical Issues

A file called *ex5_files.zip* is found in the course website. The zip contains the test files, on which your program will be tested. If any problem was discovered during the test (i.e., your program returned a different value than expected in one of the method calls, or an exception was thrown and not caught), an error message will be printed to the screen. Otherwise, if you passed successfully all of the tests, “*Perfect!*” will be printed. As in previous exercises, after you submit your exercise, we will run your code against all the input files, and you will receive an email with the tests you failed on (if any). Error messages will contain some information about the problem. For example, you may receive the following error: "runTests[17](Ex5Tester): /cs/course/2013/oop/scripts/ex5/tests/020.txt (# simple right-left rotation after delete): line number 7: expected:<true> but was:<false>"

This means that your program failed the test in the file *020.txt*, when trying to execute the delete operation. More specifically, line 7 of the test file tests the delete operation, and the test failed at this stage because the return value of calling *delete()* was true instead of false.

For your convenience, you may (and are advised to!) run the automatic tests from the terminal by running the following command:

```
~ oop/bin/ex5AutoTests.sh ex5.jar
```

(where *ex5.jar* is your jar file)

Good Luck!

Appendices

A Test Input File Format

We will run automatic tests on your code. In this appendix, we describe the format of these tests. Each of the automatic test is incorporated in a file of the following format:

- The first line starts with a '#' sign, and is a comment describing the test (i.e., this line is ignored).
- The second line creates an *AvlTree* object (denoted *myTree*) using one of the constructors defined in Section 4.3:
 - An empty line calls the default constructor.
 - A list of integers (e.g., *1 5 7 12 4*) calls the data constructor with this list of integers.
 - An existing Avl tree
- The next lines call the methods of the created tree. The lines are called one after the other. They can be one of the following:
 - *add number return_value*
call *myTree.add(number)*. Verify that the method's return value matches *return_value* (either *true* or *false*).
 - *delete number return_value*
call *myTree.delete(number)*. Verify that the method's return value matches *return_value* (either *true* or *false*).
 - *contains number return_value*
call *myTree.contains(number)*. Verify that the method's return value matches *return_value* (a non-negative integer or -1).
 - *size return_value*
call *myTree.size()*. Verify that the method's return value matches *return_value* (a non-negative integer).
 - *copy*
call *new myTree2(myTree)*. Verify that the copy constructor is written well by iterate over *myTree* and *myTree2* nodes and check that their values are equal.
 - *minNode number return_value*
call *AvlTree.findMinNodes(number)*. Verify that the method's return value matches *return_value*.

A.1 Examples

A.1.1

Consider the following example test file:

```
# a simple creation of a tree and adding the number 5

add 5 true
size 1
```

In this example, the first line is a comment describing the test. The second line is empty, which means calling the empty constructor. The third line calls *myTree.add(5)*, and then verifies that the command succeeded (by verifying that the return value of the *add* method is *true*). The last line verifies that the tree size is now 1 (after adding a single element).

A.1.2

Consider another example:

```
# calling the data constructor with the values (1 2 3) , deleting 3, and then searching for 4  
1 2 3  
delete 3 true  
contains 4 -1
```

The first line is, again, a comment describing the test. The second line is a list of numbers, which means calling the data constructor with the integer array (1 2 3). The third line calls *myTree.delete(3)*, and then verifies that the command succeeded. The last line calls *myTree.contains(5)* and verifies that 5 is found in the tree (return value should be -1).

A.2 Some Clarifications

Please note the following:

- Your code passes a given test only if your methods return the correct values in **each** of the method calls.
- A test file may include numerous calls to *add*, *delete*, *contains* and *size*. However, only a single tree is constructed in every test.
- The file that reads the input files and follows their instruction, as well as the test files themselves, are supplied by the course staff.